





# Dataset Description

369  
Volumes






155 Slices per  
Volume



⋮



4 Image  
variants per  
slice  
(Native, T1,  
T2, T2-Flair)  
+  
1 image mask  
per slice



Total number of volumes: **369**

Total number of slices: **57195**

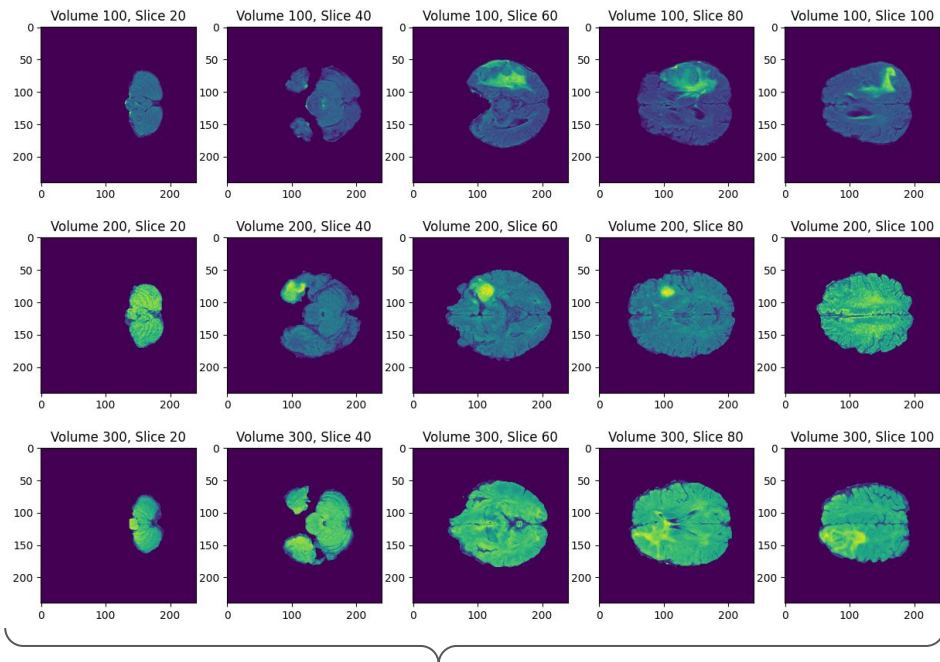
Volumes with no positive slices: **0 (0.0)**

Volumes with at least one positive slice: **369 (1.0)**


Negative slices: **32773 (0.5730046332721391)**

Positive slices: **24422 (0.4269953667278608)**

# Sample Images

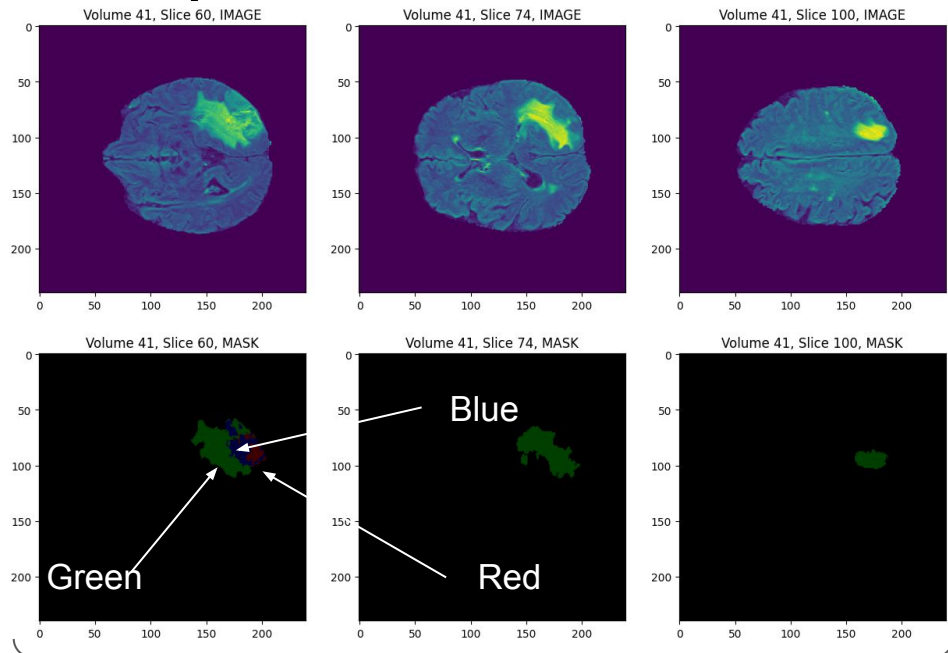


Slices of similar number correspond to same parts of the brain



Different variants highlight different structures for the same slice

# Sample Masks




Mask labels have each pixel as 4 possible values (identified by green, red, blue and black) above, which indicate (ET) GD-enhancing tumor, (ED) peritumoral edema, (NCR) necrotic and non-enhancing tumor core, and normal

# Data Exploration

Q: How many positive slices does each volume have?

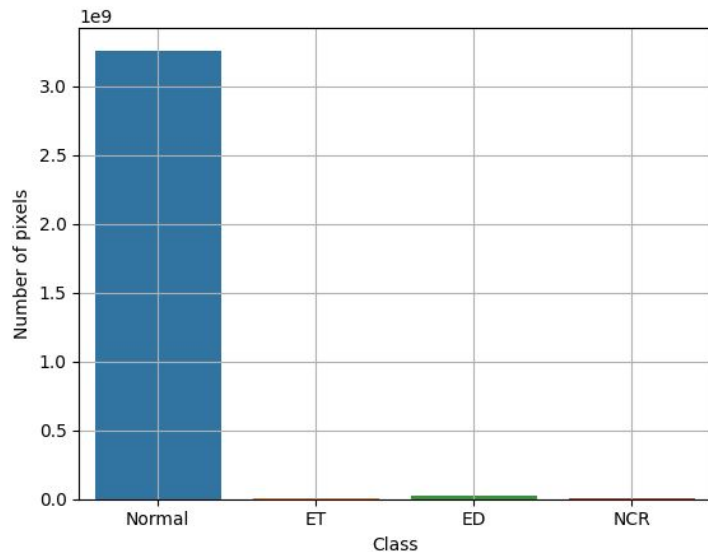
A: Each volume has between ~20 and ~100 positive slices. Most volumes have around ~60 to ~80 positive slices



# Data Exploration


Q: Is the dataset balanced with respect to the different mask classes?

A: No. The vast majority of pixels is “Normal”, probably techniques are needed to deal with this.



Q: Ignoring the “Normal” class, is the dataset balanced?

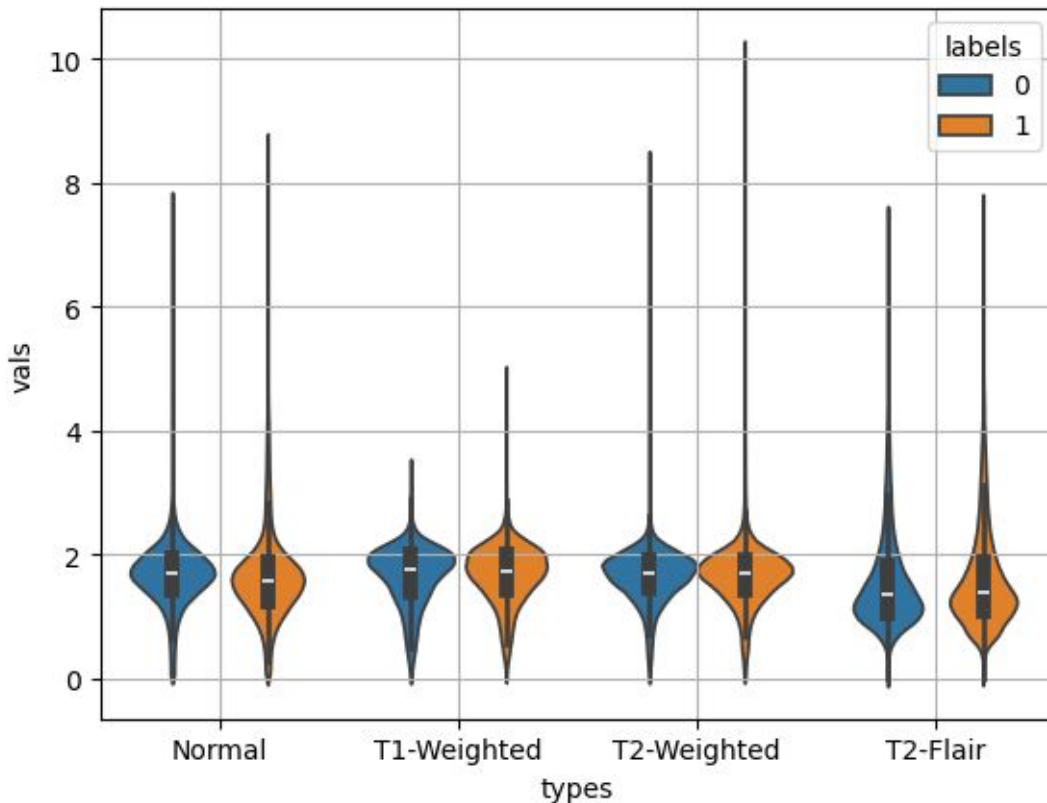
A: Not perfectly balanced, but the number of pixels of each class is comfortably within the same order of magnitude.




# Data Exploration

Q: Is the distribution of image pixel values measurably different between positive vs negative slices?


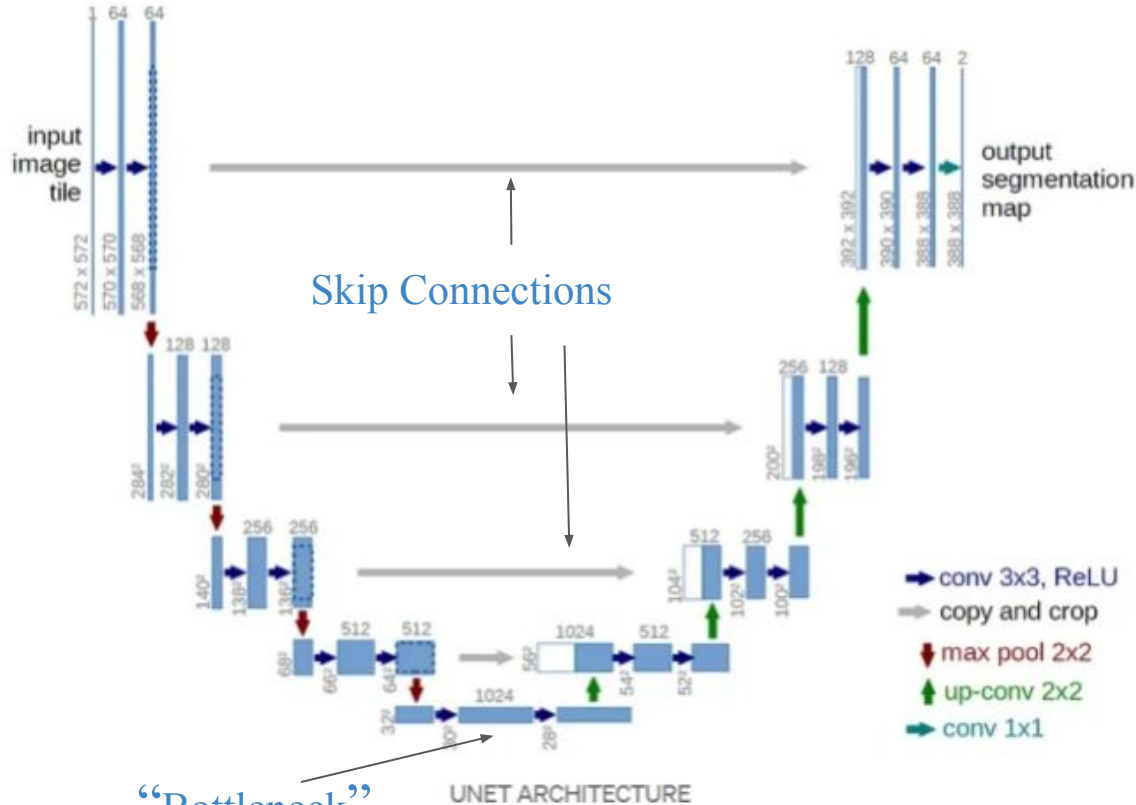
A: Distribution of pixel values is quite different between the different image variants, but the difference between positive and negative images is extremely subtle at best. In all 4 variants though, the positive images seem to have higher pixel brightness upper bounds.



# ML technique: U-Net Architecture




Encoder  
(Downsampling)



Decoder  
(Upsampling)

# U-Net Architecture


## Encoder



## Downsampling:

- **Convolutional Layers:** extract and learn hierarchical features from the 3D data
- **Active Functions:** ReLU
- **Pooling Layers:** put between convolutional layers, reduce the dimensions of the volume and downsample the data
- **Increasing Feature Channels:** number of feature channels doubles after max pooling

## Decoder




## Upsampling:

- **Upsampling Layers:** increase the spatial dimensions of the feature maps
- **Convolutional Layers:** process the merged feature maps; refine the upsampled features
- **Active Functions:** ReLU
- **Decreasing Feature Channels:** number of feature channels halves after each upsampling step; transitioning from a deep feature representation to the final segmented output




# U-Net Architecture

## Skip Connections



- Concatenation between symmetrical stages of encoder/decoder
- Retain high-resolution features and spatial information



## Bottleneck




- A narrow and deep layer that connects the encoder and the decoder
- Spatial information from the input image is compressed into a more abstract feature representation.

# U-Net Architecture

## Activation Function: rectified linear unit ReLU

- Introduces non-linearity into the model that helps the NN to learn more complex data.
- Simple and computationally efficient.
- Sets all negative values in the output to zero, while leaving the positive values as-is.



## Loss Function: pixel-wise binary cross-entropy

- Calculates the cross-entropy between the predicted probability distribution and the true probability distribution for each pixel in the image.
- Penalizes misclassification of foreground pixels, more heavily than background pixels.

$$E = \sum_{\mathbf{x} \in \Omega} w(\mathbf{x}) \log(p_{\ell(\mathbf{x})}(\mathbf{x}))$$

# Potential Training Strategy

- Stochastic gradient descent with momentum and weight decay.
- Momentum adds a fraction of the previous update to the current update in order to speed up the convergence of the optimization algorithm.
- Weight decay adds a penalty term to the loss function to have the model use smaller weights.